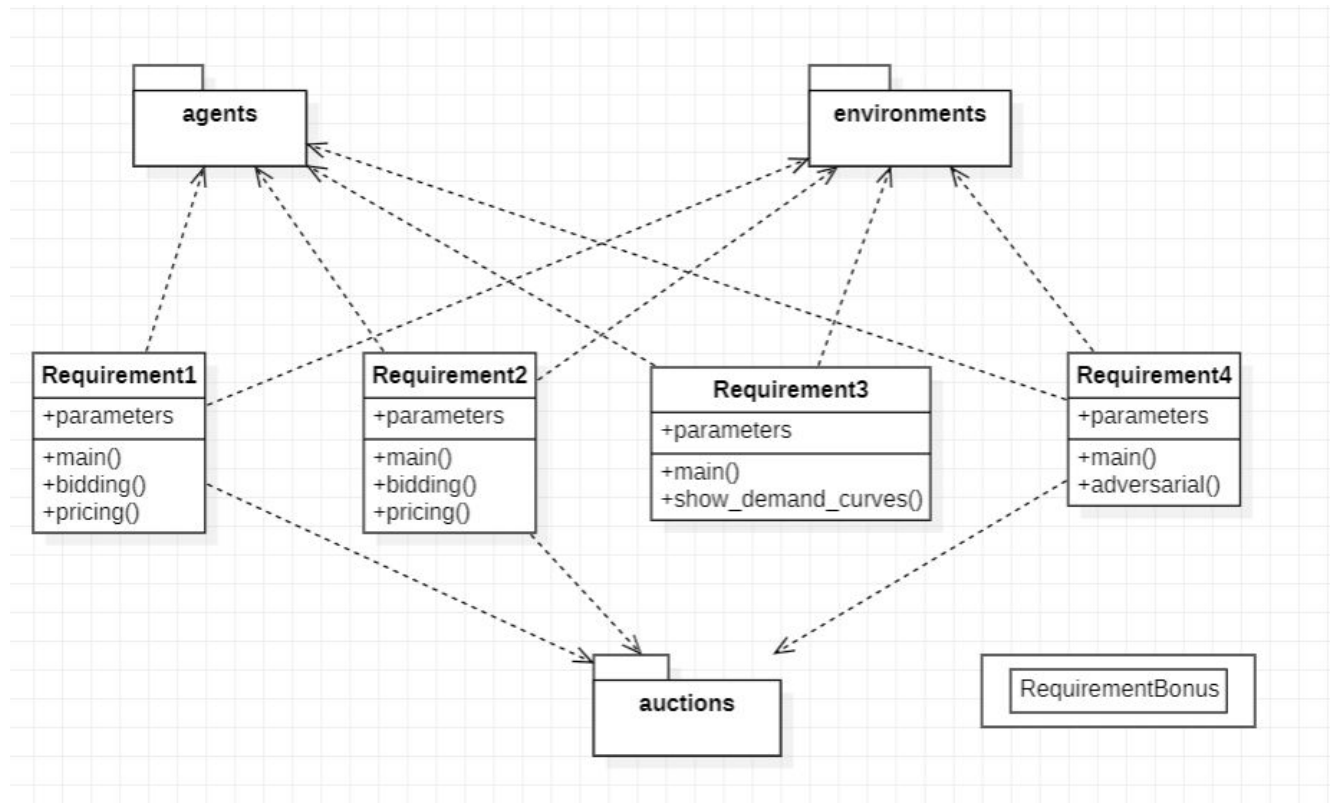


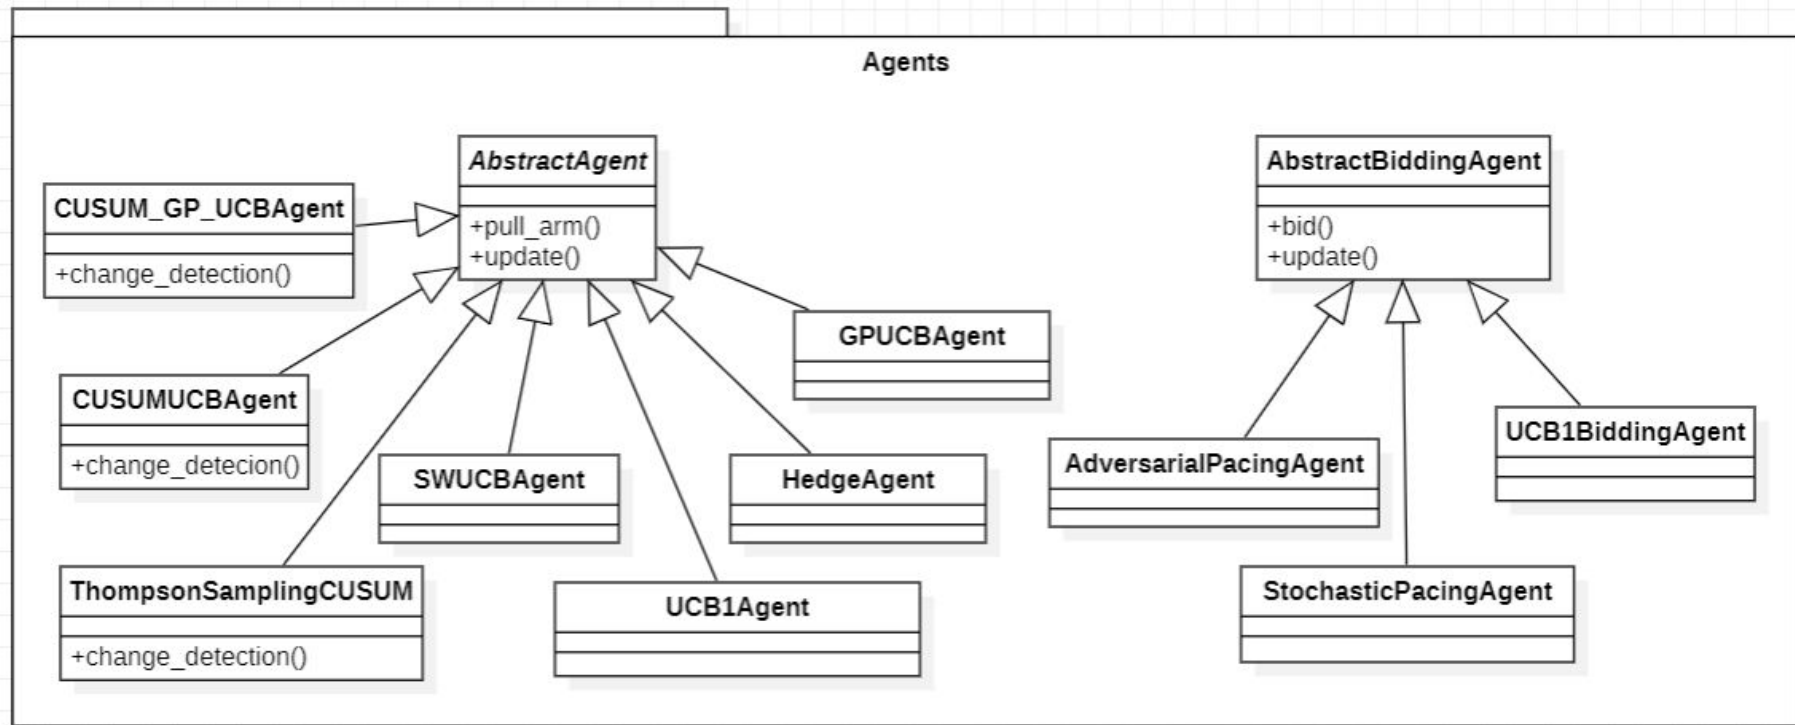
OLA24

Barda Luca
Grillo Niccolò
Franzè Lorenzo

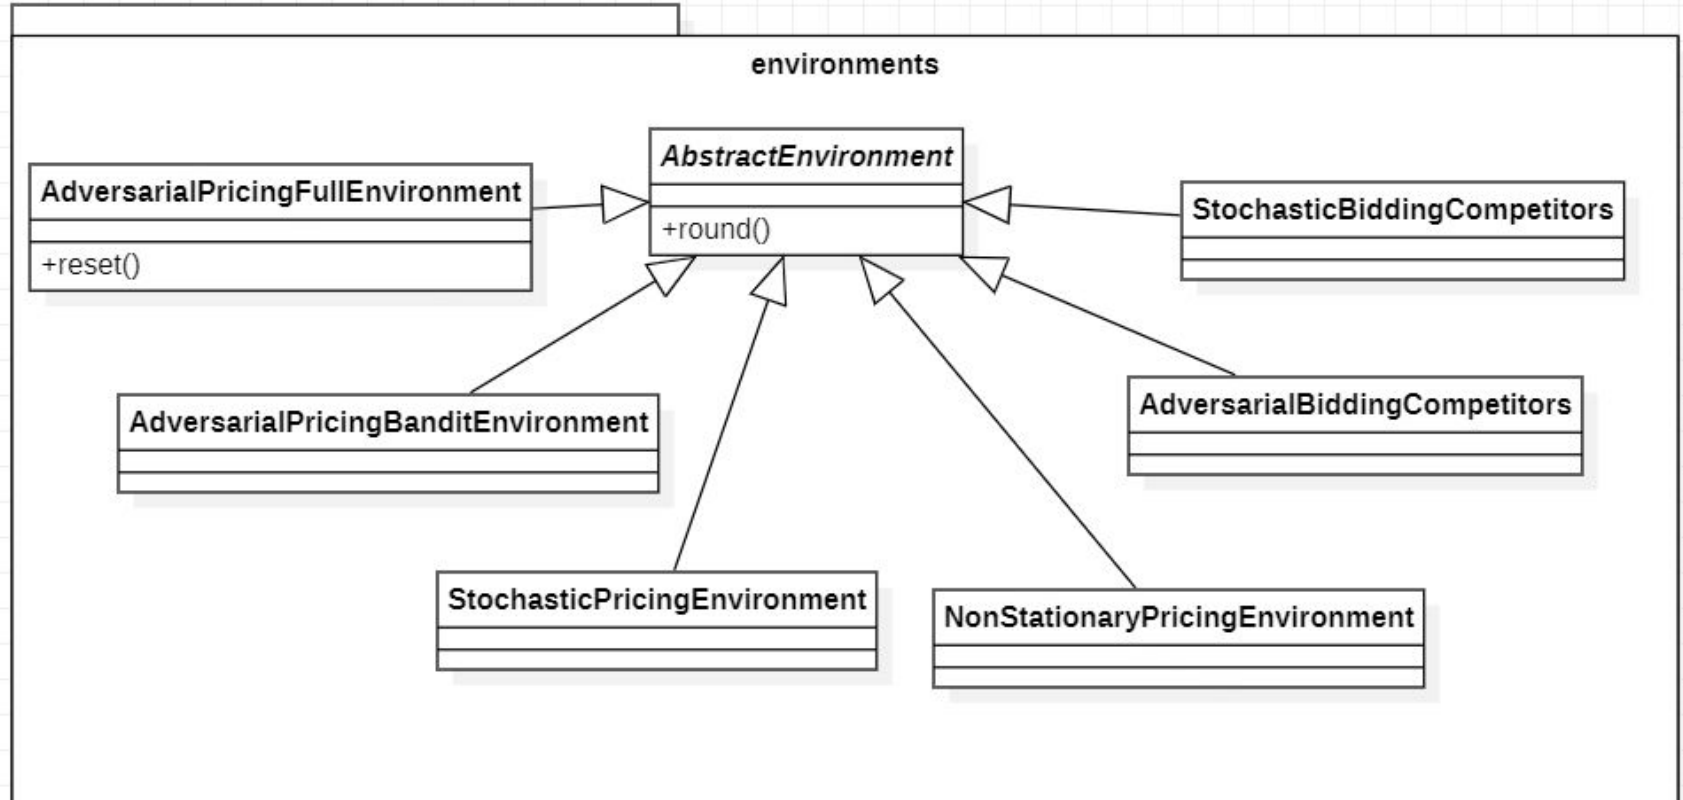
Project structure



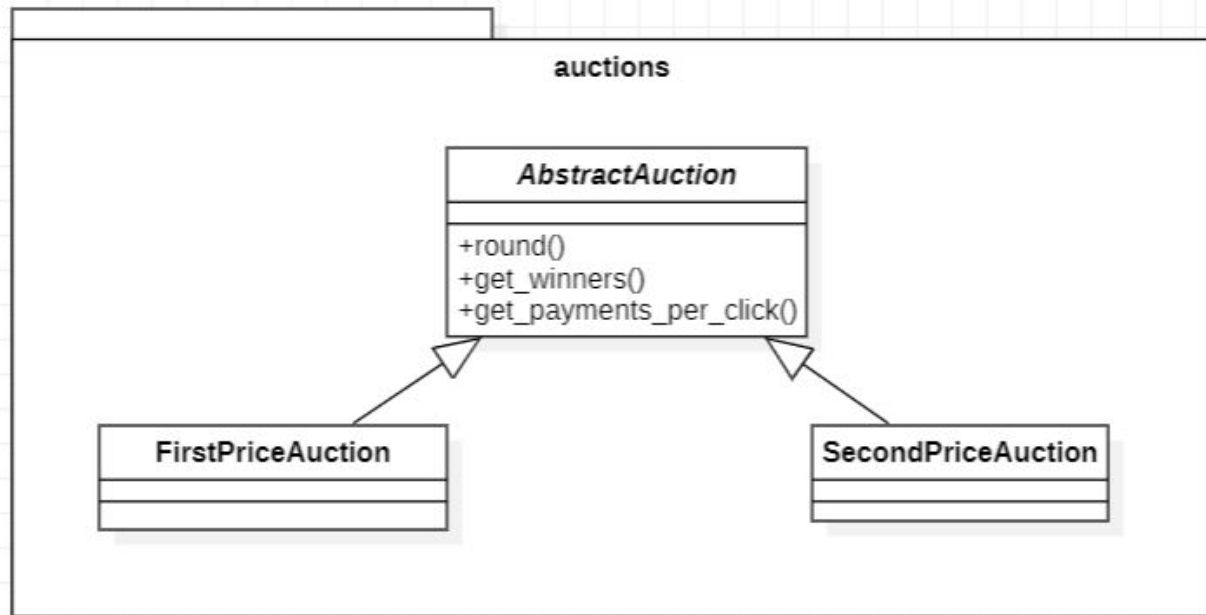
Agents package



Environments package



Auctions package



Requirement 1

```
--run_type ['main', 'bidding', 'pricing']
--bidder_type ['UCB', 'pacing']
--auctions_per_day
--num_days
--n_iters
--num_competitors
--budget
--valuation
--ctr
--theta
--item_cost
--num_buyers
```

What's left?

- demand curve
- opponent bids
- opponent ctrs

agents:

- UCB1BiddingAgent (bidding)
- StochasticPacingAgent (bidding)
- GPUCBAgent (pricing)

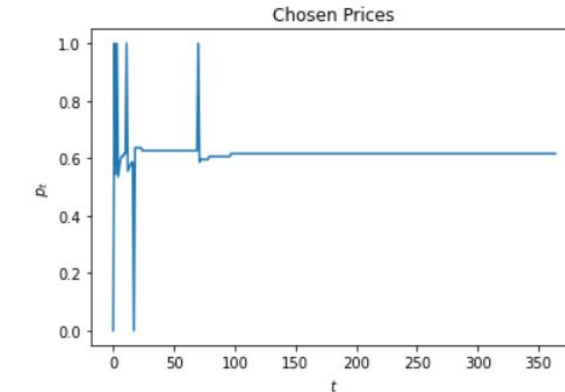
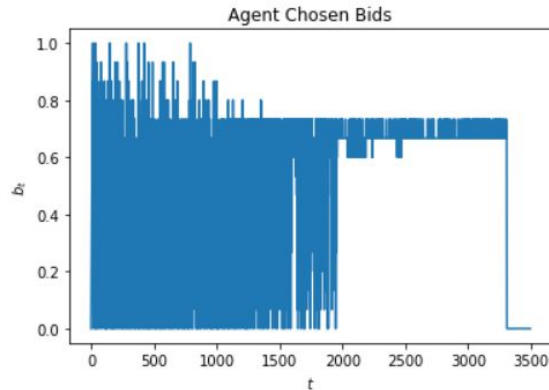

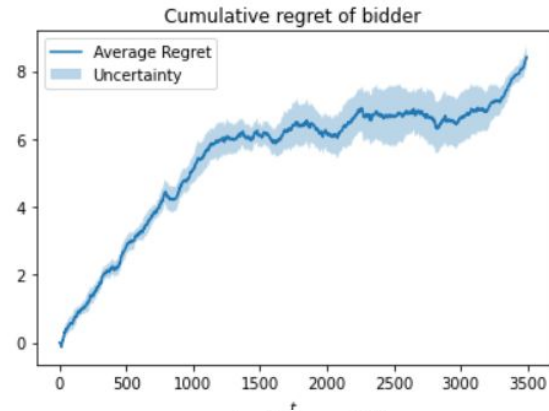

environments:

- StochasticPricingEnvironment (pricing)
- StochasticBiddingCompetitors (bidding)


auctions:

- SecondPriceAuction

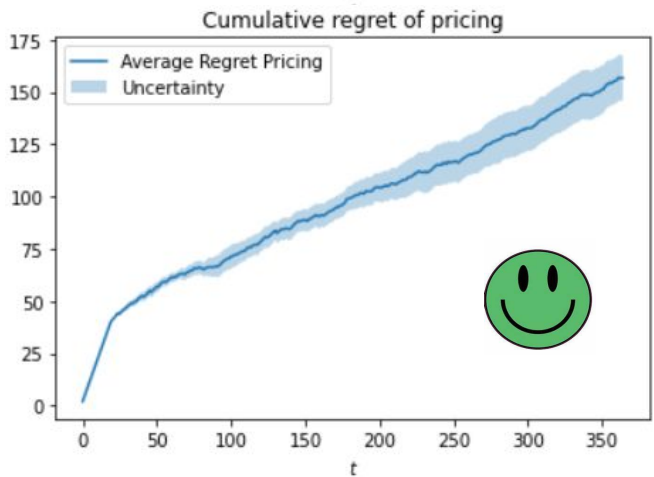


Results bidding and pricing



Results main



increase
auctions_per_day



Requirement 2 - Adversarial Full Feedback

```
--run_type ['main', 'bidding', 'pricing']  
--auctions_per_day  
--num_days  
--n_iters  
--num_competitors  
--budget  
--valuation  
--ctr  
--theta  
--item_cost  
--num_buyers
```

agents:

- AdversarialPacingAgent (bidding)
- HedgeAgent (pricing)

environments:

- AdversarialBiddingCompetitors (bidding)
- AdversarialPricingEnvironment (pricing)

auctions:

- FirstPriceAuction

Requirement 2 - clairvoyant

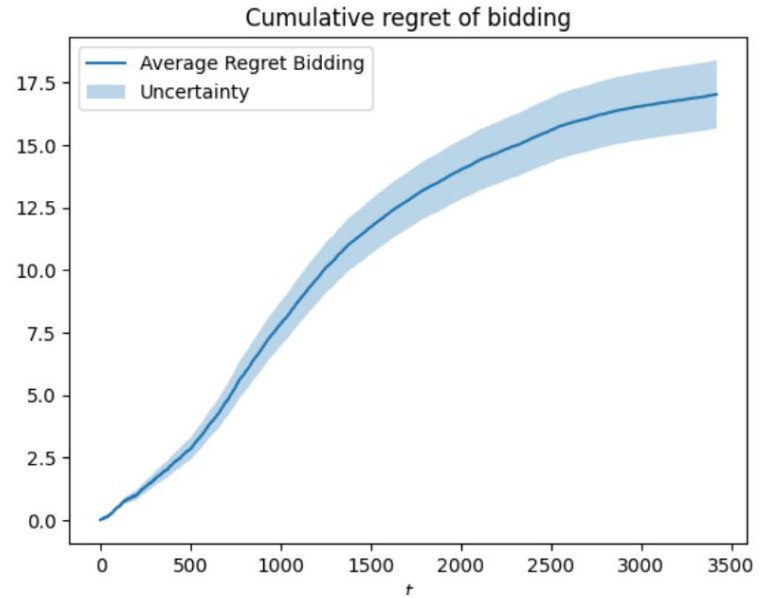
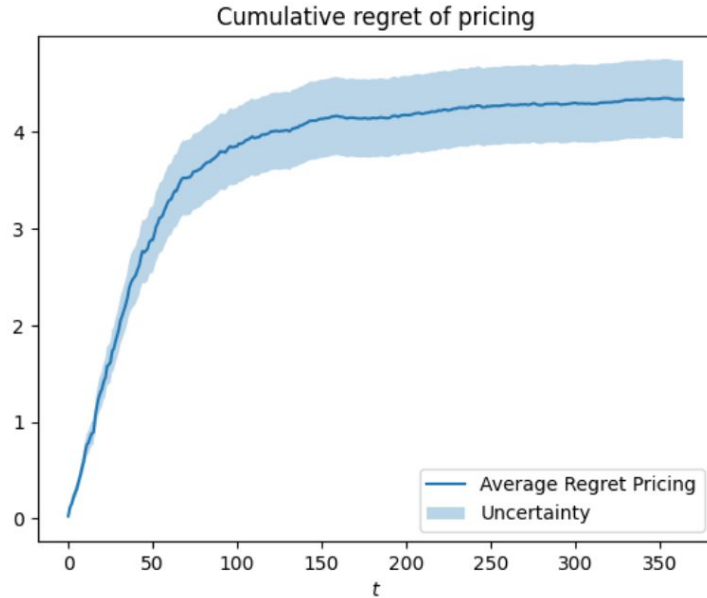
- We used the **adversarial clairvoyant** for both the pricing and bidding:
- For the bidding part is built with the following strategy

Algorithm 1 Clairvoyant Utility in Adversarial Auctions

Ensure: Maximize utility under budget constraint

```
1:  $max\_utility \leftarrow -\infty$ 
2: for each  $bid$  in  $discr\_bids$  do
3:    $c \leftarrow 0$  ▷ Total money spent
4:    $utility \leftarrow 0$ 
5:   for each auction  $t = 1$  to  $n\_auctions$  do
6:     if  $c < budget$  then
7:        $all\_bids[idx\_agent, t] \leftarrow bid$ 
8:        $winner, \_ \leftarrow auction\_agent.get\_winners(all\_bids[:, t])$ 
9:       if  $winner == idx\_agent$  then
10:         $utility \leftarrow utility + (my\_valuation - bid)$ 
11:         $c \leftarrow c + bid$ 
12:      end if
13:    else
14:      break
15:    end if
16:  end for
17:  if  $utility > max\_utility$  then
18:     $max\_utility \leftarrow utility$ 
19:  end if
20: end for
21: return  $max\_utility$ 
```

Requirement 2 - Results



*(competing bids are generated with rand. unif while the demand, for pricing the conv. prob is parametrized through a theta param randomized randomly.)

```
%run req2.py --num days 365 --auctions per day 10 --n iters 100  
--num competitors 10 --my valuation 0.8 -budget 100 --run type  
'main'
```

*(The competitors clickthrough-rates are randomized with a uniform, changing seed at each iteration.)

Requirement 3

Non-stationary demand curves

```
num_buyers = 100
T_pricing = 50000
intervals = 5
T_interval = 10000
cost = 10
max_price = 40
K = T_interval**(0.33) = 20
```

Agents:

- **UCB1 agent**
- **Sliding Window UCB agent**
- **GP-UCB CUSUM agent**
- **CUSUM UCB agent**
- **Thompson Sampling CUSUM agent**

D_1 :

$$D_1(\text{price}) = \max\left(0, 1 - \frac{\text{price}}{30}\right)$$

D_2 :

$$D_2(\text{price}) = \max\left(0, 1 - \frac{\text{price}}{60}\right)$$

D_3 :

$$D_3(\text{price}) = \max\left(0, \exp\left(-\frac{(\text{price} - 10)^2}{25}\right)\right)$$

D_4 :

$$D_4(\text{price}) = \max\left(0, \frac{1}{2\sqrt{0.05 \cdot \text{price} - 0.3}} - 0.5\right)$$

D_5 :

$$D_5(\text{price}) = \max\left(0, 1 - \left(2.4 \cdot \frac{\text{price} - 10}{30} - 2.8 \cdot \left(\frac{\text{price} - 10}{30}\right)^2 + 1.4 \cdot \left(\frac{\text{price} - 10}{30}\right)^3\right)\right)$$

Environment:

NonStationaryPricingEnvironment

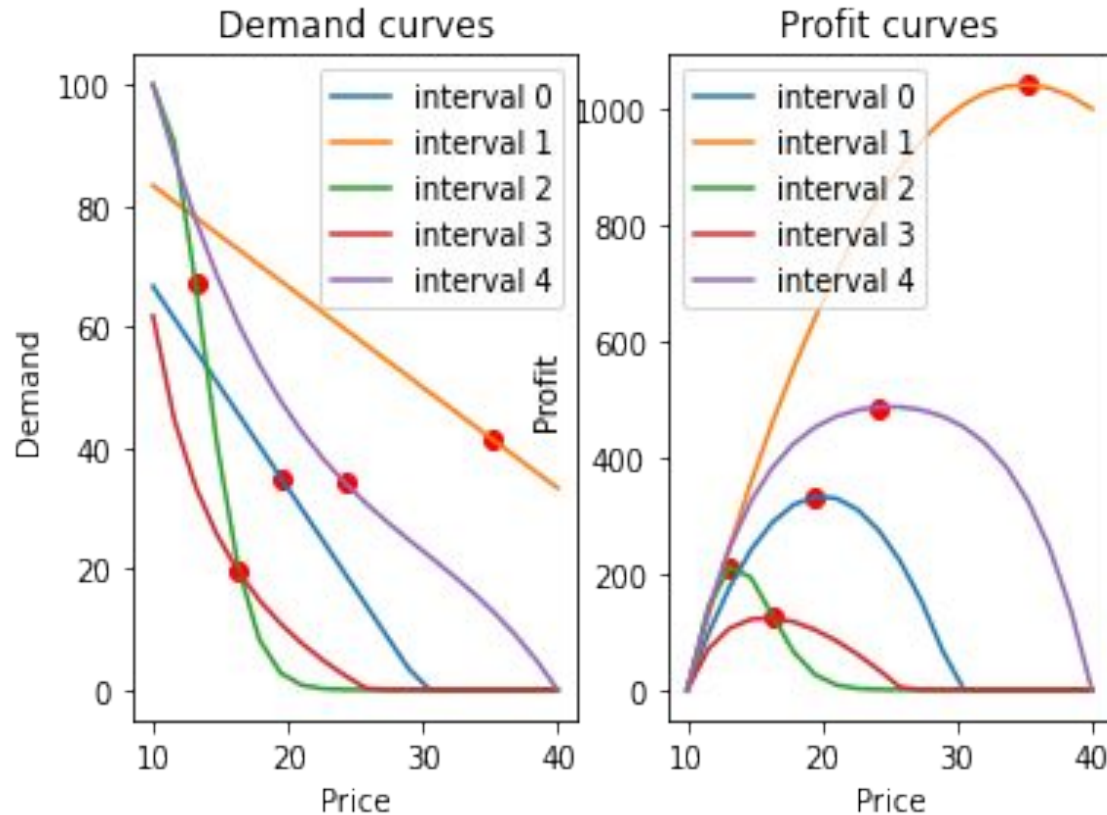
- **Noise:** at each round Binomial distribution to get the numbers of buyers out of all the buyers, according to probability given from the Demand curve
- **Non-stationary:** 4 change points in which the demand curve changes

Clairvoyant:

best prices for each interval: [19.473684210526315, 35.26315789473684, 13.157894736842106, 16.315789473684212, 24.210526315789473]

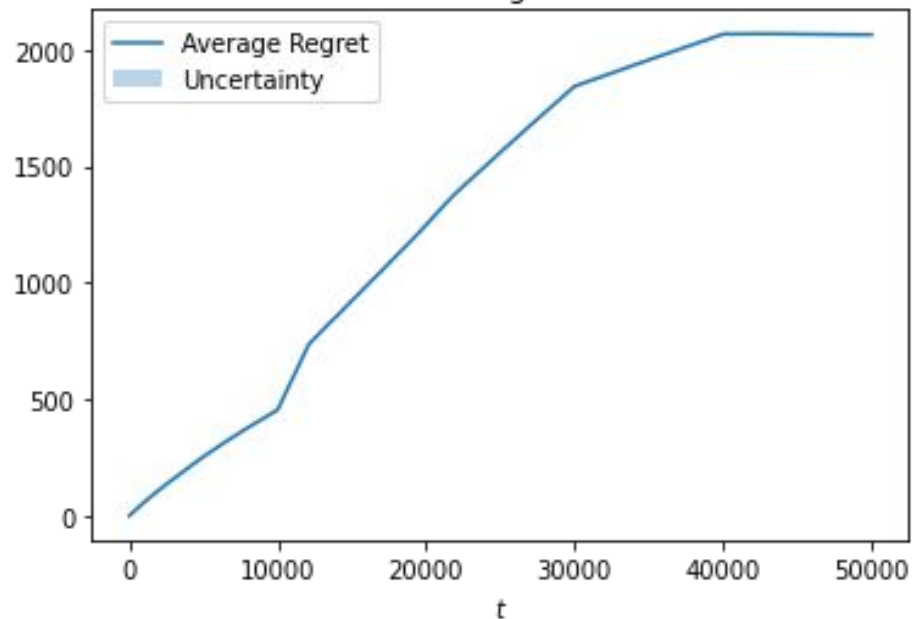
best profit for each interval: [332.4099723 1041.55124654 211.91469302 123.91553435 486.87471704]

maximum profit: 3000.0

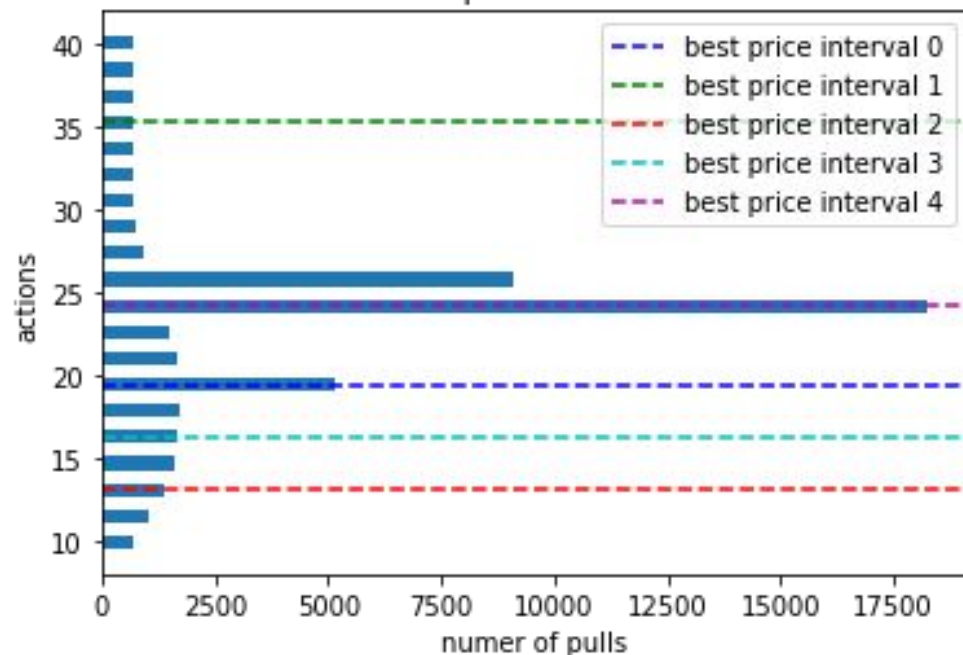


UCB 1 Agent

cumulative regret of UCB1



Number of pulls for each action




Sliding Window UCB Agent

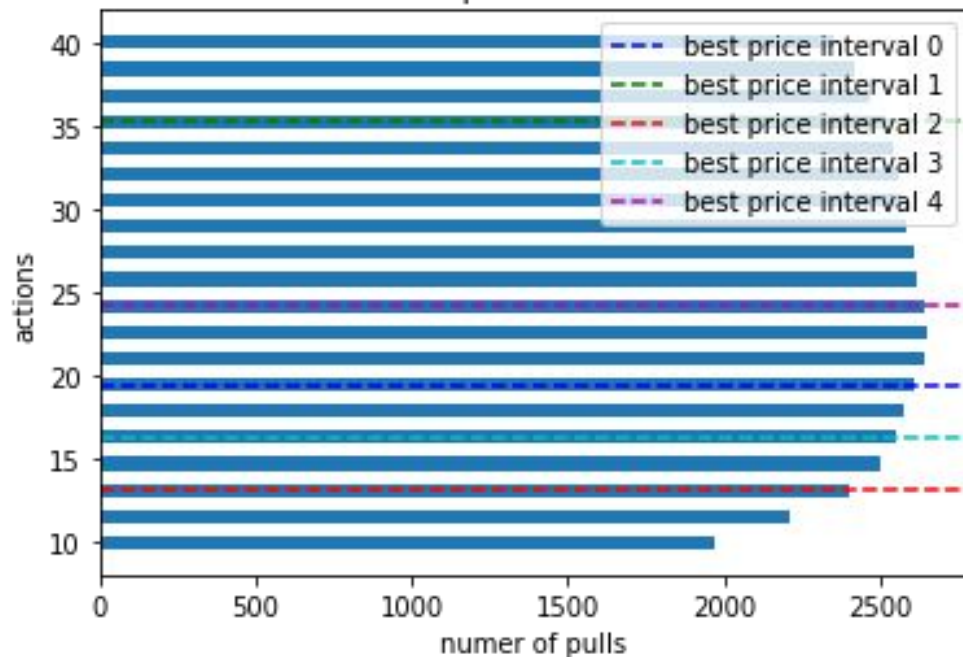
Window size: 735

$$W = \left\lfloor 2 \cdot \sqrt{\frac{T_{\text{pricing}} \cdot \log(T_{\text{pricing}})}{\text{intervals} - 1}} \right\rfloor$$

cumulative regret of UCB sliding window



Number of pulls for each action



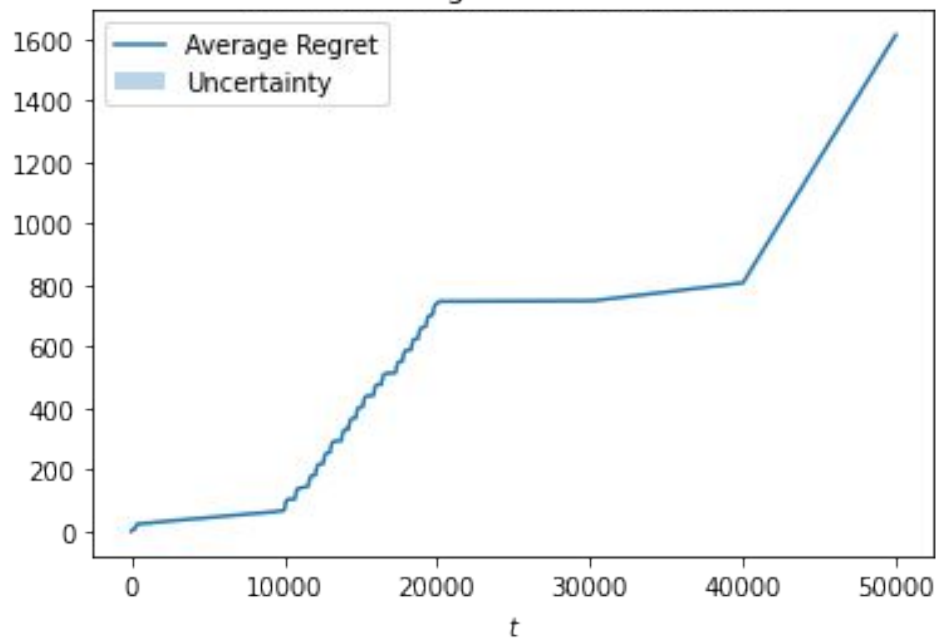
GP-UCB CUSUM Agent

Used tuned parameters

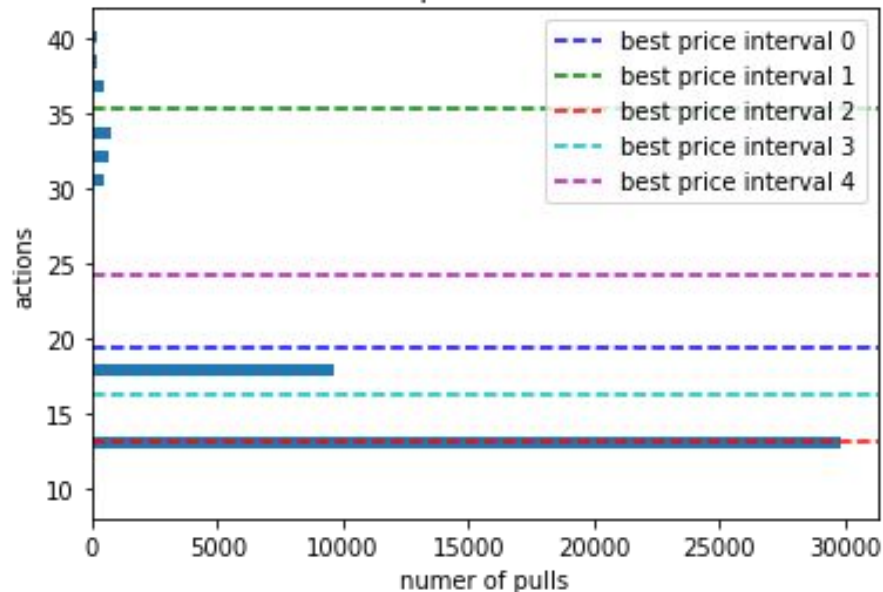
$$M = 2 \cdot \left\lceil \log \left(\frac{T_{\text{pricing}}}{\text{intervals} - 1} \right) \right\rceil$$

$$h = 240, \quad \epsilon = 130$$

cumulative regret of GP-UCB-CUSUM



Number of pulls for each action



Considerations:

- Parameters tuned in order to get the best results however not all the change points were detected
- Tradeoff between number of changes (ex. in the second interval many wrong changes are detected)
- Variant of the UCB - GP algorithm since in order to reduce the execution time if the same action is performed for more than 40 times with a tolerance of $1e-8$ than the GP isn't updated anymore
- For each change detected all the actions are restored: the whole GP is restarted from scratch

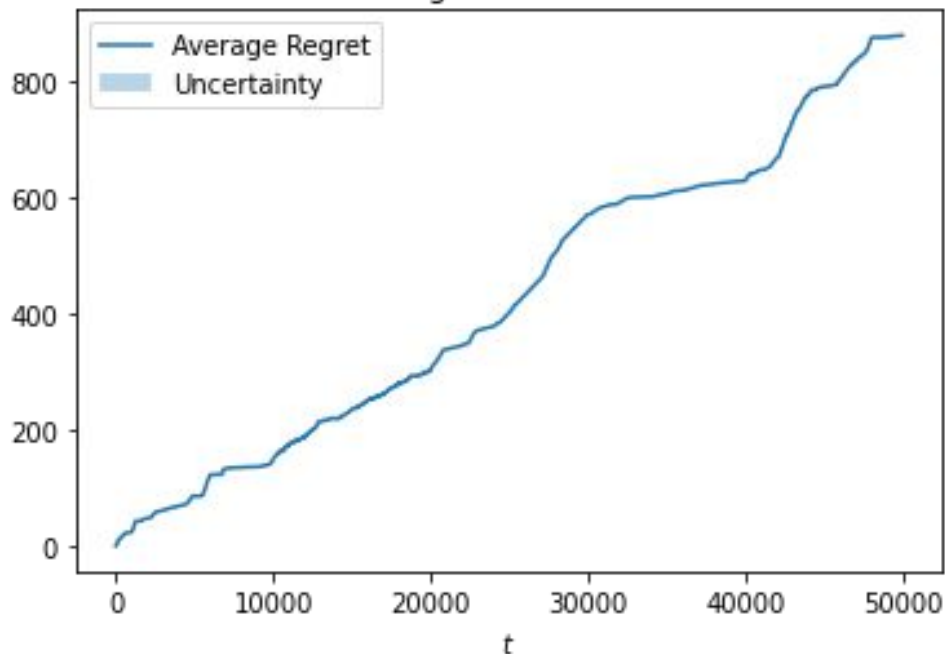
CUSUM UCB Agent

$$h = 2 \cdot \log \left(\frac{T_{\text{pricing}}}{\text{intervals} - 1} \right)$$

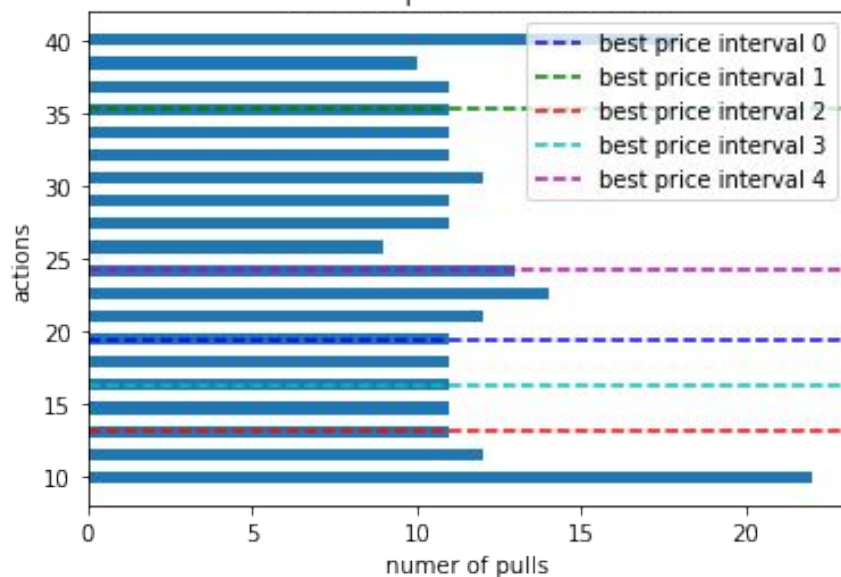
$$\alpha = \sqrt{\frac{(\text{intervals} - 1) \cdot \log \left(\frac{T_{\text{pricing}}}{\text{intervals} - 1} \right)}{T_{\text{pricing}}}}$$

$$M = \left\lfloor \log \left(\frac{T_{\text{pricing}}}{\text{intervals}} \right) - 1 \right\rfloor$$

cumulative regret of UCB with CUSUM



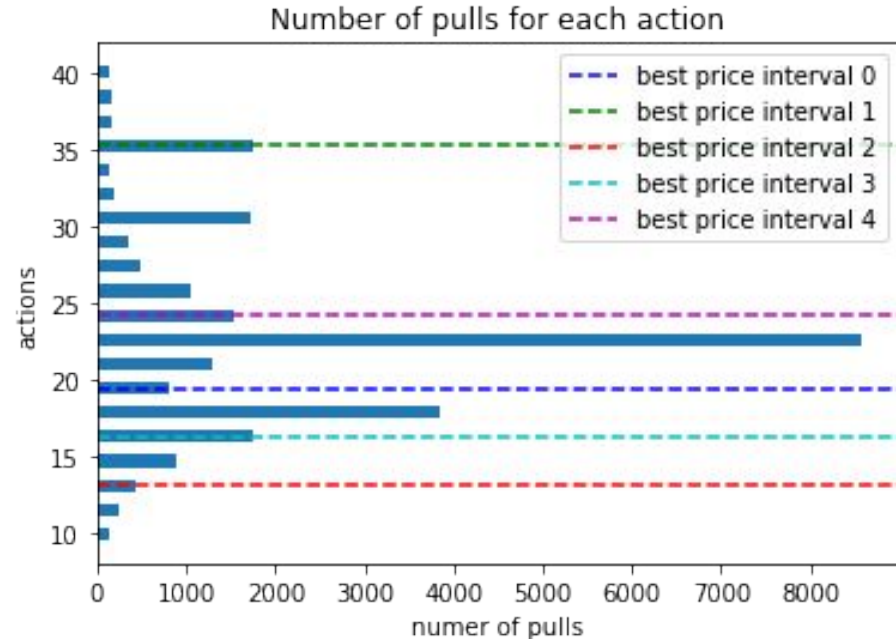

Number of pulls for each action



Thompson Sampling CUSUM Agent

h and M as described by theory

cumulative regret of UCB with CUSUM vs Thompson Sampling with CUSUM



Requirement 4 - Three bidders' show-down

```
--scenario ['solo, 'adversarial]
```

```
--num auctions
```

```
--n_iters
```

```
--num auctions
```

```
--num participants
```

```
--budget
```

```
--valuation
```

```
--my_ctrs
```

```
--theta
```

```
--eta
```

```
--num buyers
```

agents:

- StochasticPacingAgent
- AdversarialPacingAgent
- UCB1BiddingAgent

environments:

- None
or
- AdversarialBiddingCompetitors



auctions:

- FirstPriceAuction

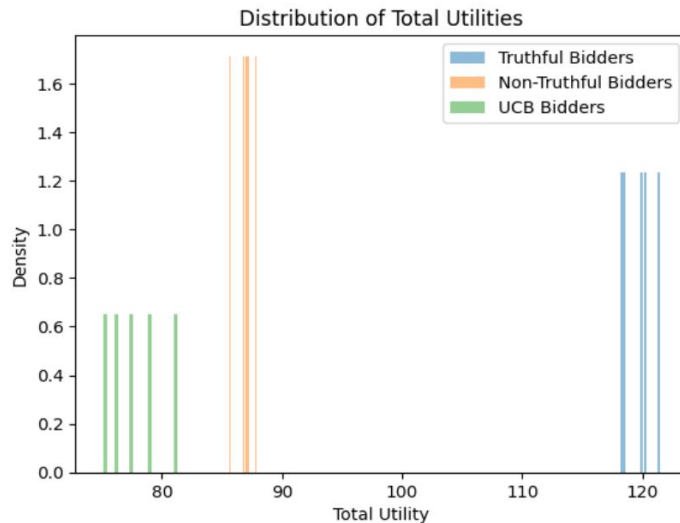
Two Scenarios:

1. num_participants is made entirely of bidders of the 3 types each with num_participants//3 members.
2. There are 3 agents (1 for each type) and (num_participants - 3) bidders

Requirement 4.1 - solo





- Here the **Regret** is computed by **averaging** across the regrets (adversarial clairvoyant) of all the bidders of each type. → **Unstable** metric.





```
%run req4.py --num auctions 3000 --budget 300 --n iters 5
```

```
--num participants 30 --eta 0.01 --seed 1 --scenario solo --valuation 0.7
```

Requirement 4.2 - *adversarial*




- Here we consider the 3 Regrets (adversarial clairvoyant) considering only the bids of the adversarial bidders. → More stable metric.




```
%run req4.py --num auctions 5000 --budget 200 --n iters 30 --num participants 100 --eta  
0.01 --seed 1 --scenario adversarial --valuation 0.8 --my ctrs 0.8,0.8,0.8
```

Requirement 4.2 - *adversarial*

- Set η to the value from theory (higher than before)



- Raise the ctr of the non-truthful bidder from 0.8 to 0.85



Requirement bonus

Demand curves



$$D_1(p_1, p_2) = \max \left(0, \left(1 - \beta_1 \times \left(\frac{p_1}{\text{max_price}} \right) + \gamma_1 \times \left(\frac{p_2}{\text{max_price}} \right) \right) \right)$$

$$D_2(p_1, p_2) = \max \left(0, \left(1 - \beta_2 \times \left(\frac{p_2}{\text{max_price}} \right) + \gamma_2 \times \left(\frac{p_1}{\text{max_price}} \right) \right) \right)$$

$$D_{\text{both}}(p_1, p_2) = \max \left(0, \left(1 - \text{both_factor} \times \left(\frac{p_1}{\text{max_price}} \right) - \text{both_factor} \times \left(\frac{p_2}{\text{max_price}} \right) \right) \right)$$

$$\begin{aligned} d_{\text{total}}(p_1, p_2) = & \left(\frac{\text{num_buyers_market1}}{\text{num_total_buyers}} \right) \cdot D_1(p_1, p_2) \\ & + \left(\frac{\text{num_buyers_market2}}{\text{num_total_buyers}} \right) \cdot D_2(p_1, p_2) \\ & + \left(\frac{\text{num_buyers_both}}{\text{num_total_buyers}} \right) \cdot \text{curve_both}(p_1, p_2) \end{aligned}$$

Users interested in buying only one item



Demand curve modelization for two items
market:

item 1 cost : 10

item 2 cost : 15


max price : 80

num buyers total : 220

Furthermore we know 20 are interested in
buying both the products

Time duration : 1000

Discretization of the price space : 100



Profit curve and Environment

$$\begin{aligned}\text{profit_curve} = & \text{num_buyers_market1} \cdot D_1(p_1, p_2) \cdot (p_1 - \text{cost1}) \\ & + \text{num_buyers_market2} \cdot D_2(p_1, p_2) \cdot (p_2 - \text{cost2}) \\ & + \text{num_buyers_both} \cdot D_{\text{both}}(p_1, p_2) \cdot (p_1 + p_2 - \text{cost1} - \text{cost2})\end{aligned}$$

$$d_{t_1} \sim \text{Bin}(\text{num_buyers_market1}, D_1(p_1, p_2))$$


$$d_{t_2} \sim \text{Bin}(\text{num_buyers_market2}, D_2(p_1, p_2))$$

$$d_{t_{\text{both}}} \sim \text{Bin}(\text{num_buyers_both}, D_{\text{both}}(p_1, p_2))$$

$$d_t = d_{t_1} + d_{t_2} + d_{t_{\text{both}}}$$


$$r_t = (p_1 - \text{cost1}) \cdot d_{t_1} + (p_2 - \text{cost2}) \cdot d_{t_2} + (p_1 + p_2 - \text{cost1} - \text{cost2}) \cdot d_{t_{\text{both}}}$$

Total Profit Curve





best price for item 1: 60.90909090909091
 best price for item 2: 80.0
 total profit : 5158.290289256198
 Expected number of users that would buy for
 each population type at the found prices:
 Market 1: 12.41477272727273
 Market 2: 46.53409090909091
 Both markets: 12.954545454545453

Total Profit Curve in the noisy environment



Agent: GP UCB Multidimensional



Best price for item1 found: 61.61616161616162

Best price for item2 found: 80.0